

IVAN SHABANOV

ivan@ishbnv.dev · +374 77 719 518 · linkedin.com/in/ishbnv · github.com/ishbnv · ishbnv.dev

Yerevan, Armenia (GMT+4) · Ready to relocate to Tallinn, Estonia · Remote until relocation, full EU overlap

11 September 2026

Bolt — People & Workplace Team

Re: AI / Agent Engineer, Tallinn, Estonia

Dear Hiring Team,

I am applying for the AI / Agent Engineer role within the People and Workplace team at Bolt. My work combines backend systems with three years of shipping LLM agents, model cascades, and safe tool interfaces, with a steady focus on evaluation, auditability, and guardrails.

At SODA, our lead-generation agency relied on human chat managers who took about eight minutes to respond to requests. To automate this process safely, I led the development of an LLM conversational assistant and ran a three-month controlled comparison across three modes: human staff, humans assisted by AI, and AI acting alone. On identical metrics, the autonomous assistant won every month and cut response time to under thirty seconds. We shifted the channel entirely to the AI assistant, tripling monthly revenue from about \$47K to about \$141K. Later, I introduced a confidence-based model cascade where a lightweight model handles routine requests and escalates only on low confidence, cutting API costs tenfold without degrading output quality.

Safety boundaries and auditability have been central to my work since. At Alfa-Dengi, I introduced AI coding agents across our engineering team with strict boundaries, preventing agents from touching database migrations or financial logic. I built custom MCP servers exposing database schemas and task trackers, and added an agentic pre-review stage that reduced module delivery from about two hours to roughly forty minutes. At Appbooster, I modeled an asynchronous validation lifecycle behind a moderation queue with anti-fraud rules, ensuring 100% of payout operations remained fully auditable across hundreds of daily transactions.

My production services run on TypeScript, Node.js, PostgreSQL, Redis, Apache Kafka, and Kubernetes on AWS, connecting agents to Anthropic and OpenAI APIs. I am comfortable working in Python codebases and building evaluation suites to keep non-deterministic systems dependable.

I live in Yerevan, Armenia, GMT+4, and am fully ready to relocate to Tallinn, with immediate full European workday overlap.

When automating People processes involving sensitive employee data, where do your human-in-the-loop checkpoints add the most value: verifying policy classifications or approving multi-step action plans before tool execution?

Kind regards,

Ivan Shabanov