

IVAN SHABANOV

ivan@ishbnv.dev · +374 77 719 518 · linkedin.com/in/ishbnv · github.com/ishbnv · ishbnv.dev

Yerevan, Armenia (GMT+4) · Ready to relocate to Tallinn, Estonia · Remote until relocation, full EU overlap

11 September 2026

Bolt — People & Workplace Team

Re: Senior AI / Agent Engineer, Tallinn, Estonia

Dear Hiring Team,

I am applying for the Senior AI / Agent Engineer role on the People and Workplace team at Bolt. My work combines distributed systems engineering with three years of architecting production LLM agents, establishing safe tool contracts, and setting team standards for agentic evaluation and human oversight.

At SODA, our lead-generation agency relied on human chat managers who took about eight minutes to respond to requests. To automate this process safely, I led the development of an LLM conversational assistant and ran a three-month controlled comparison across three modes: human staff, humans assisted by AI, and AI acting alone. On identical metrics, the autonomous assistant won every month and cut response time to under thirty seconds. We shifted the channel entirely to the AI assistant, tripling monthly revenue from about \$47K to about \$141K. Later, I introduced a confidence-based model cascade where a lightweight model handles routine requests and escalates only on low confidence, cutting API costs tenfold without degrading output quality.

Setting standards and safe operational boundaries has been central to my leadership. At Alfa-Dengi, I established our team's AI-agent engineering practices: defining strict guardrails that prevented agents from modifying migrations or financial logic, and building custom MCP servers to expose database schemas and task trackers through structured tool contracts. I also added an agentic pre-review stage that reduced module delivery from about two hours to roughly forty minutes. At Appbooster, I modeled an asynchronous validation lifecycle behind a moderation queue with anti-fraud rules, ensuring 100% of sensitive payout operations remained fully auditable across hundreds of daily transactions.

My production services run on TypeScript, Node.js, PostgreSQL, Redis, Apache Kafka, and Kubernetes on AWS, connecting agents to Anthropic and OpenAI APIs. I know when deterministic automation is superior to an LLM, and I am comfortable working in Python codebases.

I live in Yerevan, Armenia, GMT+4, and am fully prepared to relocate to Tallinn, with immediate full European workday overlap.

As you build out the People automation roadmap, where do you see the biggest architectural challenge: orchestrating multi-system rollbacks across disparate HR tools, or establishing evaluation ground truths for subjective employee queries?

Kind regards,

Ivan Shabanov