

IVAN SHABANOV

ivan@ishbnv.dev · +374 77 719 518 · linkedin.com/in/ishbnv · github.com/ishbnv · ishbnv.dev
Yerevan, Armenia (GMT+4) · Ready to relocate to London, UK · Remote until relocation, full EU overlap

11 September 2026

Flo Health — AI Platform Hiring Team

Re: Senior AI/ML Engineer (London, UK)

Dear Hiring Team,

I am applying for the Senior AI Engineer position on the AI Platform team at Flo. Over the past three years, I have built and operated production LLM systems where end-user perception of latency, tone, and reliability directly impacted revenue and engagement. Flo's mission to deliver medically grounded, personalized GenAI features at an 80-million-user scale matches my experience designing model cascades, custom tool contracts, and automated notification workflows.

At SODA, I led the development of an automated AI lead-processing assistant where response time and answer quality were critical to user trust. To validate the system, I ran a three-month controlled test comparing human operators, human-plus-AI assistance, and an autonomous AI flow. The autonomous AI workflow consistently outperformed human baselines, slashing first-response latency from roughly eight minutes to under thirty seconds and tripling monthly channel revenue from about \$47K to about \$141K. To resolve the inference cost and latency trade-off, I designed and shipped a confidence-based model cascade: a fast, lightweight model handles routine inquiries, escalating to a heavy model only when certainty drops below threshold. This halved latency and cut LLM API expenditures tenfold without degrading answer quality.

At Alfa-Dengi, I bridged complex backend logic with intuitive user journeys. I built a visual drag-and-drop campaign canvas in React and NestJS that allowed non-technical marketers to construct multi-step automated notification workflows with frequency caps, generating about \$80K a month. To accelerate team-wide engineering velocity, I authored project conventions for Claude Code, built custom MCP servers exposing database schemas and documentation to agents, and introduced an agentic pre-review stage that dropped CRUD delivery time from roughly two hours to about forty minutes.

My core engineering stack spans TypeScript, Node.js, React, REST and WebSocket APIs, PostgreSQL, Redis, Kafka, Docker, and AWS, and I am comfortable working with Python across LLM pipelines. I design for structured outputs, clear tool schemas, and systematic evaluation against ground truth.

I am currently based in Yerevan, Armenia (GMT+4) and fully prepared to relocate to London for this hybrid role, while ready to begin collaborating remotely under full UK timezone overlap during the visa and onboarding process.

With Flo serving millions of active users daily, where does the AI Platform team encounter the sharpest trade-off today: maintaining sub-second streaming latency during peak morning engagement, or calibrating clinical safety evaluation gates across personalized user journeys?

Kind regards,

Ivan Shabanov