

IVAN SHABANOV

Senior AI Engineer • LLM Systems, AI Agents, MCP • TypeScript / Node.js • Product Engineering

Yerevan, Armenia (GMT+4) • Ready to relocate to London, UK • Remote until relocation, full EU overlap
ivan@ishbnv.dev • +374 77 719 518 • linkedin.com/in/ishbnv • github.com/ishbnv • ishbnv.dev

SUMMARY

AI Engineer with 4+ years of production software engineering and 3 years shipping LLM systems: an automated AI assistant that tripled monthly channel revenue to ~\$141K, a confidence-based model cascade that cut LLM API costs 10×, and custom MCP servers adopted across an engineering team. Combines product sensibility with systems execution: prompt and tool contract design, latency and cost optimization, structured outputs, evaluation against human baselines, and integration into high-load applications. Ready to relocate to London.

EXPERIENCE

Alfa-Dengi via GIDFINANCE — Senior Full-Stack Engineer • Tech Lead

Sep 2024 – Present • 2 yrs • Remote

Consumer-lending marketplace with 9K+ new leads a day. Cross-functional product team; contract engagement via GIDFINANCE.

- Introduced AI-agent workflows across the team: authored **project rules for Claude Code** (architectural guardrails, conventions), built **custom MCP servers** exposing database schemas and documentation to agents, and added an agentic pre-review stage; typical module delivery dropped from ~2 h to ~40 min.
- Enabled marketers to launch automated push campaigns **with zero engineering involvement** (~\$80K/month generated) by building a visual drag-and-drop campaign canvas on React and NestJS with normalized state and per-status frequency caps.
- Launched a credit-report product that **turned profitable in its first week** (net profit +40%, EPC doubled from ~\$1.9 to ~\$3.8) by leading a 3-engineer team through technical design, user flows, API contracts, and a Node.js backend querying a **6M-row** scoring dataset.
- Eliminated notification loss** under peak traffic bursts by moving notification delivery and domain-event processing to **Apache Kafka**, establishing reliable, replayable event streaming.
- Cut load time of a 100K-row lead-scoring table (9K+ new leads/day) from **8 s to under 1 s** by reworking the data-access path: server-side pagination, PostgreSQL query and index optimisation, and client-side list virtualization.

Appbooster — Full-Stack Engineer (contract)

Nov 2023 – Aug 2024 • 10 mos • Remote

Mobile-app marketing platform: ~100 employees, 10,000+ apps promoted for clients in 26 countries.

- Grew the task-and-rewards product to **1,000 daily active users in 3 months** by owning the full cycle: React user dashboard for paid tasks and proof submission, chatbot re-engagement notifications, and a live verification queue for admins.
- Made **100% of payouts auditable** and screened out duplicate and faked proofs by modelling the full task lifecycle from pick-up to payout behind a moderation queue with anti-fraud rules, processing **hundreds of completions a day**.

SODA — Full-Stack Engineer

Aug 2022 – Oct 2023 • 1 yr 3 mos • Remote

Lead-generation agency for residential developers and law firms (~60 employees). Owned the client-server side: APIs, database, interfaces.

- Tripled monthly channel revenue from ~\$47K to ~\$141K** and cut lead response time from ~8 min to **under 30 s** by leading an AI lead-processing assistant across a 3-month controlled comparison (humans / humans + AI / AI only) that moved the channel fully to AI.
- Cut LLM API costs 10×** on own initiative by prototyping and shipping a model cascade: a fast, cheap model handles routine leads and escalates to a stronger model only on low confidence.
- Lifted conversion ~1.5× on SPA sites receiving paid ad traffic by raising Lighthouse Performance from **58 to 87** through Core Web Vitals work, asset loading, and layout stability.

OPEN-SOURCE PROJECT

vk-ai-bot-platform (MIT) — multi-tenant AI chatbot platform: OpenRouter LLM API, per-community prompt versioning, token-aware context windows, multi-step re-engagement workflows on BullMQ, event deduplication, admin analytics. TypeScript, Fastify, Drizzle ORM, PostgreSQL, Redis, React, Docker Compose, GitHub Actions.

SKILLS

AI Engineering & LLM Systems: LLM integration (Anthropic API, OpenAI API, OpenRouter), AI agents, agentic workflows, MCP (Model Context Protocol) custom servers, prompt engineering, structured outputs, model cascading and routing, cost/latency optimisation, LLM evaluation against baselines, Claude Code

Languages & Stack: TypeScript, JavaScript (ES6+), Node.js, SQL, ready to work with Python in project stack

Backend & APIs: NestJS, Express, Fastify, REST APIs, GraphQL, WebSocket, microservices architecture, system design, Apache Kafka, BullMQ, OWASP Top 10

Frontend & Product UX: React, Next.js, Redux Toolkit, visual workflow builders, Core Web Vitals, Lighthouse, interaction design, accessibility

Data & Infrastructure: PostgreSQL, Redis, Docker, Kubernetes, AWS (EC2, S3, networking, deploy), GitLab CI/CD, GitHub Actions, OpenTelemetry, Grafana

Testing & Process: Vitest, Jest, Supertest, Playwright (E2E), A/B testing, code review, technical documentation, mentoring

EDUCATION & LANGUAGES

Southern Federal University — Applied Informatics, 2022–2024

English — professional working proficiency (daily in a distributed team) • Russian — native