

IVAN SHABANOV

ivan@ishbnv.dev · +374 77 719 518 · linkedin.com/in/ishbnv · github.com/ishbnv · ishbnv.dev
Yerevan, Armenia (GMT+4) · Remote · B2B contract or EOR · Full overlap with EU & IST hours

11 September 2026

ServiceTitan — Site Reliability & Infrastructure Engineering Team

Re: Senior Site Reliability Engineer (Job ID JR115793)

Dear Hiring Team,

I am applying for the Senior Site Reliability Engineer role on the Site Reliability & Infrastructure Engineering team at ServiceTitan. My work focuses on operating Kubernetes clusters in production, building observability with OpenTelemetry and Grafana, and developing AI-assisted SRE workflows with Claude Code and Model Context Protocol.

At Alfa-Dengi, a consumer-lending marketplace handling over 9,000 new leads daily, I designed our deployment and observability infrastructure from scratch. I built automated GitLab CI/CD pipelines deploying directly into a 5-node self-managed Kubernetes cluster, paired with OpenTelemetry distributed tracing and Grafana alerting. This telemetry pipeline enabled our team to spot query regressions and latency degradation before users reported issues. When promotional bursts previously resulted in dropped notifications, I moved notification delivery and domain events to Apache Kafka, creating an asynchronous, replayable event stream that eliminated message drops entirely under peak traffic. On our primary lead-scoring table of 100,000 rows, I profiled execution plans and tuned PostgreSQL queries and indexes, reducing response latency from eight seconds down to under one second.

ServiceTitan's focus on AI-assisted engineering with Claude Code and MCP-based agents mirrors what I have built over the past two years. At Alfa-Dengi, I authored project conventions and reusable skills for Claude Code, defining strict boundaries around production migrations and financial logic. I built custom MCP servers that expose database schemas, task trackers, and internal documentation directly to agents, and established an agentic pre-review stage before human code review. This automation reduced typical CRUD module delivery with tests from roughly two hours to about forty minutes.

Across production environments, I manage AWS infrastructure including EC2 instances, S3 storage, and VPC networking security groups. At SODA, I decomposed a legacy monolith into containerized microservices running on Linux with automated pipelines, and I use Redis and BullMQ for caching and queueing.

I work remotely from Yerevan, Armenia (GMT+4), which sits just one hour and a half behind Bengaluru (IST) with near-total workday overlap, and I am available for B2B contract or EOR engagement.

As ServiceTitan scales MCP-based agents for live incident investigation, how do your SRE teams currently establish verification guardrails to ensure autonomous remediation agents never execute unsafe state mutations during high-severity outages?

Kind regards,

Ivan Shabanov